

Kiliaan Vanvoorden

Riemst, Belgium · bakerstreetbandit@zohomail.eu · 0032497774869 · github.com/BoozeLee

Self-taught engineer who builds the verification layer for AI systems: the gates, tests and CI that decide whether a system is right, and that fail loudly when it is not. Every figure below comes from a command you can re-run. Looking for a first role where writing the evaluation harness is the job.

Projects

harness — **contract-gated AI engineering pipeline** · public · **AGPL-3.0**

- <https://github.com/BoozeLee/harness>
- Ship an 11-subcommand pipeline (`scan` → `task` → `verify` → `pr`) that blocks protected-path writes and secret reads before an agent's edit lands. 88 tests passing.
- Measure the harness against the agent it replaces: 13 graded runs against a hidden acceptance suite, 11 passed.
- Report the unflattering number too — scope violations ran 3 raw versus 4 harnessed; the contract was amended after the runner flagged them.
- Install with pip from the repo (1.1.0, wheel builds via `uv build`); `scan --fail-under 70` exits 1 below the floor, a blocked guard write exits 2.

elohim — **checksum-pinned gate for numerical claims** · public · MIT

- <https://github.com/BoozeLee/elohim>
- Pin every published number to a checksum: 81 facts classified, 107 checksum-pinned, 17 declared exemptions, 0 unclassified.
- Catch instrument tampering before trusting a result: 6 of 6 seeded mutations caught, each emitting `verdict FAIL` and a `PIN DRIFT` line naming the altered file.
- Ship 7 measurement skills behind one command (`python3 tests/test_all.py` → `ALL_SKILLS_PASS`); CI and CodeQL green.
- 22,196 non-blank tracked lines (Python 21,980) across 43 commits (measured at `a047ccd`).

terminal221b — **bounded local-first coding CLI with a Rust TUI** · public · **AGPL-3.0**

- <https://github.com/BoozeLee/terminal221b>
- Bound agent context by path, not by a token budget; 445 tests (385 TypeScript, 60 Rust) cover `path-guard`, `boundary-drift` and `scope`.
- Redact credentials out of stored transcripts before they reach disk, with a test that fails if a secret survives.
- Ship a TypeScript CLI with a `ratatui` terminal UI alongside an Expo client. 16,210 non-blank tracked lines, 55 commits.

repotruth — **CI that reads the repository, not just the diff** · public · **AGPL-3.0**

- <https://github.com/Bakery-street-project/galacticfederation>
- Verify every webhook: accept a signed push, answer a replay with an idempotent 200, reject an unsigned push with 400 and queue nothing, reject a tampered body holding an otherwise-valid signature.
- Report `truncated: false` and its own input limits instead of silently truncating an audit above them.
- Expose one repository knowledge graph three ways — MCP server, CLI, fleet service. 147 tests across 32 suites.

Also public

- **mcp-regression-lab** — a GitHub Action that diffs an MCP server's tool contract across releases, so a renamed or narrowed tool fails CI instead of failing in production. 29 tests.
- **mycroft** — fixed-scope AI diagnostics that end in a written Minimum Viable Action.

Private, available to show on request

- **capo** — 349 tests across five packages as one gate, holding a plan-schema contract before execution.
- **superbrain** — 298 passing tests; a model may propose a numeric claim but cannot commit one.
- **beehive-studio** — 54,550 non-blank tracked lines. Its test suite is not declared in a form I have verified, so I make no test claim for it.

Skills

Use regularly — Python · TypeScript · Rust · shell · Git · GitHub Actions · `pytest` · `node --test` · `vitest` · `cargo test` · npm/pnpm workspaces · MCP tool contracts · agent skill manifests · claim-binding evaluation harnesses · tamper-detection fixtures · HMAC webhook verification

Not shipped — Docker · Kubernetes · Go

Education

Self-taught. No degree, no certifications.

Availability

Based in Riemst, Belgium. Open to remote roles, and to employment or contract.